

Cool — Pitch Pack v2

For: Pranaav Shrinaath S (Keni) · Founder, Cool **Positioning:** The PQC-enabled, decentralized, ZK-proofed runtime for AI. **Linux for AI.**

Positioning thesis

What Linux did for servers, Cool does for AI:

- **Open and neutral** — not owned by OpenAI, Anthropic, AWS, or any nation-state
- **Decentralized** — workloads run across an independent node network, not on one company's cloud
- **Verifiable** — every inference produces a cryptographic attestation + ZK proof anchored on-chain
- **Post-quantum secure by default** — all channels use NIST-standardized PQC; built for the next 30 years, not the last 10
- **Portable** — any model, any framework, deploys on the runtime

You're not building an OpenAI competitor. You're building the **substrate** that all AI eventually runs on — including OpenAI's. This positioning is what makes you category-defining instead of category-derivative.

How to use this document

Three pitches at three lengths, plus objection handling and rehearsal notes:

1. **30-second elevator** — bus, lift, conference line
2. **2-minute mid-form** — Twitter Space, coffee chat, panel
3. **5-minute full pitch** — investor first meeting, YC interview, deep founder convo

Rehearsal rule: read each one aloud 10 times before saying it to a human. Time yourself. Don't compress-speak; cut content instead.

Memorize the structure, not the words. Beat-by-beat fluency means you can answer

interruptions without losing your thread.

1. The 30-second elevator pitch

"What Linux did for servers — open, neutral, portable infrastructure that nobody controls — Cool is doing for AI.

Today, every AI workload runs on someone else's centralized cloud, on weights you can't verify, secured by cryptography that quantum computers will break within a decade.

Cool is a decentralized runtime where AI workloads execute inside hardware enclaves across an independent node network. Every inference produces a zero-knowledge proof you can verify on-chain. Every channel is post-quantum encrypted by default.

It's the substrate for AI that's sovereign, auditable, and built for the next 30 years — not the last 10."

Timing target: 30–35 seconds at natural pace **Energy:** declarative, not pitchy. You're stating a thesis, not selling a product. **The hook test:** if the listener doesn't ask one of *"how does that work?"* / *"who's it for?"* / *"who else is doing this?"* within 5 seconds — the pitch missed. Practice until the hook lands.

Variation for crypto-native audience (drop "Linux" frame, lean into the cryptography):

"AI today is centralized, unverifiable, and quantum-vulnerable. Cool fixes all three. Decentralized TEE node network, ZK proofs for every inference, post-quantum-secure channels end-to-end. The verifiable AI runtime. Solidity attestation contracts on Base L2, Rust orchestration, Phala + Nvidia CC for the TEE layer."

2. The 2-minute mid-form pitch

For: Twitter Space, panel intro, coffee with a peer founder, conference networking.

"I'm Pranaav, building Cool. The thesis: AI today runs on a stack that breaks on three axes — centralization, verifiability, and quantum security — and the consequences of those breaks are about to land all at once.

Centralization first. OpenAI, Anthropic, AWS, Google — five companies control where AI workloads run. For consumer chat that's fine. For a hospital, a court, a defense agency,

or a regulator, single-vendor lock-in at the inference layer is a structural risk nobody's built around yet.

Verifiability second. When a provider tells you GPT-5 wrote a response, you trust them on the model identity, the weights, and the output. No cryptographic proof. The EU AI Act, HIPAA's AI-decision rules, and FINRA's model-governance updates landing in 2026 and 2027 all require an audit trail that 'trust me' can't satisfy.

Quantum security third. Every TLS handshake securing AI inference today uses elliptic-curve cryptography that's harvestable now and decryptable within ten years. NIST standardized post-quantum replacements last year. Nobody's built them into the AI inference path.

Cool is the runtime that fixes all three. Decentralized — workloads run across an independent TEE node network on Phala and Nvidia Confidential Compute, not anyone's cloud. Verifiable — each inference produces a TEE attestation and a zero-knowledge proof, anchored on-chain via Solidity contracts on Base L2. Post-quantum secure — every channel between client, node, and chain is hybrid PQC-protected using NIST ML-KEM.

The framing I use: what Linux did for servers, Cool does for AI. Open, neutral, portable infrastructure that nobody controls.

I'm currently solo, sophomore-year B.Tech, came in from applied cryptography — built a post-quantum IPsec research prototype this year in C with liboqs. Looking for a co-founder with deep Rust + TEE depth, and early design partners in regulated AI."

Timing target: 2:00–2:20 **Energy curve:** start calm and analytical (the three problems), shift to declarative confidence at "Cool is the runtime that fixes all three." That's the emotional turning point.

3. The 5-minute full pitch — structured by beat

Six beats. Each has a target time, the words to anchor to, and the **job of the beat** — what the listener should be thinking by the time you stop talking.

Beat 1: The Problem (0:00 – 1:00)

Job of the beat: Listener leans in. They realize there are three problems, not one, and they're connected.

"I want to walk you through a thesis about AI infrastructure that I think most people haven't quite put together yet — but it's going to matter a lot over the next five years.

Today, every AI workload in production runs on one of five companies' centralized cloud. OpenAI, Anthropic, Google, AWS, Azure. When you call one of their APIs, you're trusting them on three things that you fundamentally cannot verify. First, that they actually ran the model they said — not a cheaper variant. Second, that they used the published weights — not a quietly fine-tuned version. Third, that they didn't modify the output between the model's generation and your screen.

For consumer chat, this is fine. Nobody cares whether their chatbot ran on GPT-4 or GPT-4-turbo. But there's a fast-growing class of use cases where this *does* matter. A radiologist using AI to support a cancer diagnosis. A bank approving a mortgage. A court accepting AI-generated evidence. A defense system identifying targets. A government deciding welfare eligibility.

In every one of those, three things are true at once. The audit trail is critical. The regulator is watching. And 'the AI said so' is not a defensible answer."

Delivery note: Slow down on the five companies. Eye contact on each of the high-stakes verticals. Pause before "three things are true at once" — that's the line that lands the gravity of the problem.

Beat 2: Why Now (1:00 – 1:50)

Job of the beat: Listener thinks "this isn't theoretical, the wave is here."

"Three forces are converging in 2026 and 2027 that make this a now-problem, not a future-problem.

One — the regulatory wave. EU AI Act has been in force for nearly a year with explicit verifiability requirements for high-risk AI systems. HIPAA enforcement is being updated to cover AI-assisted medical decisions. FINRA's new model-governance rules require financial firms to log and explain AI-driven decisions. India's DPDP Act and similar frameworks in Singapore and the UAE are catching up. By 2027, regulated industries simply won't be able to deploy AI without a cryptographic audit trail.

Two — the trust collapse. Two incidents in 2025 — one healthcare-AI vendor silently downgrading its model to cut cost, one fraud-detection vendor whose 'AI decision' turned out to be a rule-based system in a trench coat — have already shown enterprises that 'trust me' doesn't hold. The procurement question changed overnight from 'is your AI accurate?' to 'how do we verify?'

Three — the cryptographic primitives are finally production-ready. Nvidia Confidential Compute shipped at scale in 2024. Phala Network's TEE inference network went mainnet. NIST standardized ML-KEM and ML-DSA for post-quantum cryptography in August 2024. Zero-knowledge proving systems — Groth16, RISC Zero, Succinct — became economically practical for inference-sized workloads.

Eighteen months ago, this product wasn't buildable. Today, it is."

Delivery note: This beat is where investors decide whether you've done the homework on market timing. Be specific with dates and standards names. *"August 2024"* lands harder than *"recently."*

Beat 3: The Solution (1:50 – 3:00)

Job of the beat: Listener understands what you've actually built, in concrete terms, with the *Linux for AI* metaphor locked in.

"Cool is the verifiable runtime layer for AI. The framing I use: what Linux did for servers, Cool does for AI. Open, neutral, decentralized, portable infrastructure that nobody owns.

Concretely, four pillars.

One — decentralization. AI workloads on Cool don't run on one company's cloud. They run on an independent network of TEE nodes operated by node-operators we don't control. Phala for general-purpose enclaves, Nvidia Confidential Compute for GPU workloads, AWS Nitro and Intel TDX as additional trust domains. No single party can take the network down, swap a model, or surveil the inference.

Two — verifiability. Every inference produces two artifacts alongside the output. A TEE attestation, signed by the hardware itself, proving exactly which code ran on which weights with which inputs. And a zero-knowledge proof — compact, on-chain-verifiable — that lets anyone confirm the attestation without re-running the work. Both are anchored on Base L2 via Solidity contracts. The audit trail is mathematical, not contractual.

Three — post-quantum security. Every channel — client to node, node to node, node to chain — uses hybrid post-quantum cryptography combining classical elliptic curves with NIST's ML-KEM. Inference traffic harvested today cannot be decrypted when cryptographically-relevant quantum computers arrive.

Four — portability. Developers integrate by changing one line. Instead of

`openai.complete()`, they call `cool.complete()`. Same OpenAI-compatible API. Any model — open-source, proprietary, fine-tuned — deploys to the runtime. The protocol is

open; we don't lock anyone in.

The product is built in TypeScript and Rust for the orchestration layer, Solidity on Base L2 for on-chain attestation, and C for the low-level cryptographic primitives where performance matters."

Delivery note: This is the hardest beat. The four pillars must come out *clean*. Rehearse this 30 times, not 10. If you stumble on "ML-KEM" or "Base L2" you've lost the listener.

Beat 4: Why Me + Traction (3:00 – 3:50)

Job of the beat: Listener thinks "this person can actually build this."

"I'm Pranaav. Nineteen, second-year B.Tech at SRM University in Chennai, AI/ML specialization. Solo founder right now, actively recruiting a technical co-founder with deep Rust and TEE experience.

My technical background sits at the exact intersection Cool needs. Applied cryptography is the most relevant piece. I built a research prototype called ipsec-pqc-ikev2 — open-source on GitHub — that maps post-quantum cryptography, specifically NIST's ML-KEM, into the IPsec key exchange protocol. That's in C, using liboqs, aligned with NIST PQC guidance and the IETF post-quantum drafts. The work taught me how to reason about hybrid trust stacks — combining classical and post-quantum primitives — which is exactly the architecture Cool's encryption layer uses.

On the chain side, I've shipped Solidity in production. A project called RationChain uses on-chain identity verification to prevent welfare-distribution fraud — Aadhaar-verified registrations, fifty-percent fraud reduction in simulation. Move-language work on Sui for NFT loyalty systems. Rust CLIs, full-stack TypeScript with Next.js, Drizzle, Postgres — that's my daily stack.

On Cool specifically: the architecture is locked. Working integration with Phala's TEE infrastructure is in development. On-chain attestation contracts are in design. Target for first private dev users is the next 90 days, with the first verifiable-inference demo flowing end-to-end. I run an existing freelance business — QyroBay — which keeps me capitalized while I build, and I'm finalizing applications for non-dilutive Indian deep-tech grants — SISFS, TANSEED — to extend runway through MVP."

Delivery note: Be honest about being solo, being early, and being 19. Investors detect bullshit fast. The combination of "19, sophomore, MVP not shipped" + "shipped applied PQC, shipped Solidity, daily Rust + TypeScript" is a stronger profile than dressing up early as late. The honesty IS the differentiator.

Beat 5: Market + Business Model (3:50 – 4:30)

Job of the beat: Listener sees the financial shape of the opportunity.

"Market shape. The total enterprise AI inference market is roughly \$30 billion this year, growing forty percent annually. The segment where verification is moving from optional to mandated — regulated AI in healthcare, finance, legal, defense, government — is initially smaller, call it \$2 billion in 2026, but it's the part that's growing fastest and isn't optional.

Business model — consumption pricing. Per-attested-inference fees. Premium over standard inference APIs, three-to-five-x for the verification layer, with TEE compute costs passed through. Unit economics work because the customer's alternative is hiring a compliance team to manually audit AI decisions — which is an order of magnitude more expensive.

Go-to-market is bottom-up first. SDK and protocol open-source, free for non-attested inference, paid for verified inference. We build mindshare with crypto-AI developers first — the people who already care about this. Then enterprise sales into healthcare-AI vendors as the first wedge, fintech compliance as the second, legal-tech as the third.

Funding plan. Indian non-dilutive grants now — SISFS, TANSEED, Tamil Nadu Deep Tech Policy — to extend runway through MVP. Then a \$2 to 4 million seed round around month 10, anchored either by YC or an Indian deep-tech VC. The seed funds the network bootstrap, the first ten enterprise design partners, and the cryptography-engineer hires we need to harden the runtime."

Beat 6: The Ask (4:30 – 5:00)

Job of the beat: Listener knows what you want and feels invited into a real conversation.

"I'm not raising money today. What I'm raising is context — and what's most useful from you depends on who you are.

If you're a founder or engineer with Rust and TEE depth — I'm actively looking for a technical co-founder. I'd love a longer conversation.

If you're working in healthcare-AI, fintech, legal-tech, or defense — and your customers ask hard questions about AI audit trails — I'd love a twenty-minute call to test whether what I'm building solves the right problem for you.

If you're an investor — the seed opens around month 10. But I'd value a first conversation

now, partly because you can shoot holes in the thesis while there's still time to fix things, and partly because we should know each other before the round.

The single thing I'd most appreciate, whoever you are: tell me where the thesis is weakest. What would have to be true for Cool not to work? That's the question that helps me most right now."

Delivery note: End on the question. Never on a thank-you. Closing with a question transfers the conversational ball to them and signals confidence — *the pitch is done, now let's talk*.

Objection handling — the 7 questions you'll get every time

These are the most-likely-asked questions in approximate order of frequency. Prepare 30-second answers. Practice them until they're as fluent as the pitch itself.

Q1: "What stops OpenAI, Anthropic, or Google from doing this themselves?"

"Nothing stops them — and they probably will, for their own models. But three structural reasons it won't be enough. One, they have a vendor-incumbent conflict — being the prover and verifier of your own work isn't independent verification, it's marketing. Two, regulated enterprises want vendor-neutral verification — a runtime that works across OpenAI, Anthropic, Google, and open-source models — because they don't want trust-layer lock-in. Three, the regulators care about independence. The EU AI Act mentions third-party auditability for high-risk systems by name. Cool is built to be that neutral, third-party, cryptographic verifier."

Q2: "TEEs have been broken before. Why is this secure?"

"Right — Intel SGX has had multiple side-channel vulnerabilities, and we should be upfront. Two answers. First, defense in depth: we use multiple TEE backends — Phala, Nvidia CC, AWS Nitro, Intel TDX — so a vulnerability in one trust domain doesn't compromise the network. Second, the ZK proof layer adds a separate verifiable claim that doesn't depend solely on the TEE — it's a belt-and-suspenders approach. No single hardware vulnerability collapses the trust stack. And the PQC layer ensures that even if today's TEE attestation signatures get broken cryptographically a decade out, the channel encryption ahead of harvest-now-decrypt-later still holds."

Q3: "Why do you need a blockchain? Couldn't you just write attestations to a database?"

"Three reasons the chain matters. Tamper-resistance: a database operator can rewrite history; the chain can't. Public verifiability: any third party — a regulator, a customer, an auditor years later — can verify an attestation without needing access to our infrastructure. Censorship resistance: nobody can prevent a valid attestation from being recorded. For low-stakes use cases the chain is overkill — but for the regulated-AI use cases we're targeting, it's structurally necessary, and it's why we use Base L2 specifically — Ethereum-grade security with cents-per-transaction economics."

Q4: "Aren't you just a wrapper on Phala?"

"Phala is one trust domain in our network. The product isn't 'use Phala' — Phala is plumbing. Cool adds the on-chain attestation contracts, the ZK proof layer, the PQC channel encryption, the multi-TEE-backend orchestration, the OpenAI-compatible developer API, and — most importantly — the verifiability standard the runtime exposes. Phala by itself doesn't give you any of that. Saying we're a wrapper on Phala is like saying Stripe is a wrapper on the card networks. The integration layer is the product."

Q5: "How are you going to compete with hyperscalers' enterprise sales motion?"

"We aren't competing in their sales motion. The bottom-up developer adoption play comes first — open-source SDK, free for non-attested use. Developers at regulated companies bring Cool in through the side door, the way Datadog and Vercel did. By the time enterprise procurement notices, the runtime is already integrated. The enterprise sale then becomes 'turn on the paid tier,' not 'displace AWS.' We don't displace anyone — we're complementary infrastructure. That's why the Linux framing matters: Linux didn't sell against Microsoft; it ran underneath everything Microsoft sold."

Q6: "Who's the actual buyer? CISO? CTO? Compliance?"

"Initial signing buyer is usually the head of AI engineering, sometimes the CTO, occasionally the CISO depending on the regulated vertical. But the *demand signal* comes from compliance and legal, because they're the ones who can't sign off on AI-in-production today without an audit trail. So the actual sales motion is: engineering integrates the SDK, compliance discovers it and likes that the audit problem is solved, the budget owner approves the paid tier. We don't have to convince a buyer the problem

exists — they're already living it."

Q7: "Why you? You're 19 with no production-at-scale experience."

"Two-part answer. One, the technical credibility: I've shipped applied post-quantum cryptography work in C with liboqs that maps NIST PQC into IETF IPsec drafts — that's not undergrad coursework, it's the exact niche Cool's encryption layer needs. I'm pre-positioned in the cryptography subset of this. Two, the honest part: I haven't shipped production-at-scale. That's why a technical co-founder with TEE-at-scale experience is the number-one thing I'm recruiting for right now. The combination — me on cryptography and protocol design, a senior co-founder on TEE production engineering — covers the bar. I'm not pretending I have a senior team. I'm pretending I'm building one."

Rehearsal protocol — the next 5 days

Day 1 (today, after sleep)

- Read this whole doc once
- Pick whichever framing — Linux-for-AI or crypto-native — resonates more
- Record yourself doing the 30-second elevator on phone. Listen. Note one thing to improve.

Day 2

- Record the 2-minute. Listen.
- Drill objection answers Q1, Q2, Q3 out loud, three times each.

Day 3

- Record the full 5-minute. Listen back twice.
- Drill objection answers Q4, Q5 out loud.

Day 4

- Find a friend / classmate / Swetha — pitch the 5-minute to them live. Get questions. Note which beats they zoned out on.
- Drill Q6, Q7.

Day 5

- Re-record the 5-minute. Compare to Day 3. The delta tells you how much you've improved.
 - Now you're ready to pitch a real person — investor, founder, design partner — without it being your first take.
-

Final calibration note

The Linux-for-AI framing is a *positioning* tool, not a *technical* claim. You're not building an operating system. You're building a verifiable, decentralized, post-quantum-secure inference runtime that **feels like Linux's role** in the AI stack — open, neutral, portable, the thing that runs underneath the things people see.

Use the metaphor to make people *want* to understand. Then describe the actual product in precise terms — TEE attestation, ZK proofs, PQC channels, on-chain anchoring. Metaphor opens the door. Precision earns the conversation.

Don't oversell. The thesis is strong enough that you don't need to.